

My (Chiffon) Nguyen

San Francisco, CA, USA | hi@mychiffonn.com | github.com/mychiffonn | mychiffonn.com

RESEARCH INTERESTS

Current and future AI that are socially responsible and beneficial for broader world: **AI alignment** (pluralistic alignment, goal misgeneralization) and **AI control** (monitoring, scalable oversight); **multilingual and multicultural AI**; **AI x animals**; **human-AI collaboration** and interaction-centric models.

EDUCATION

Minerva University, College of Computational Sciences

Sep 2021 – May 2025

B.Sc in Computational Sciences (Machine Learning and Statistics), GPA: 3.7/4.0

San Francisco, CA, USA

- **Relevant Coursework:** Machine Learning (A), Bayesian Modeling (A), Statistical Modeling and Causal Inference (A), Optimization Methods (A), Probability and Statistics (A-)
- **Global Experience:** Seoul (South Korea), Taipei, Hyderabad (India), Buenos Aires (Argentina), Berlin (Germany)

SELECTED PUBLICATION

P=Preprint, R=Under Review, J=Journal, C=Conference, W=Workshop. Full publication list on [Google Scholar](https://scholar.google.com/citations?user=MyChiffon).

- [R1] **SEATauBench: Adapting Tool-Agent-User Evaluation Into Low-Resource Southeast Asian Languages**
My Chiffon Nguyen*, Aulia Adila*, Saksorn Ruangtanusak*, Kittiphath Leesombatwathana*, Vissuta Gunawan Lim*, Patomporn Payoungkhamdee, Samuel Cahyawijaya
Under review at EMNLP, 2026.

- [P1] [Anthropogenic Regional Adaptation in Multimodal Vision-Language Model](#)
Samuel Cahyawijaya, Peerat Limkonchotiwat, Tack Hwa Wong, et al. (including My Chiffon Nguyen)
Preprint, 2026.

- [J1] [A benchmark of expert-level academic questions to assess AI capabilities](#)
Center for AI Safety, Scale AI, HLE Contributors Consortium, et al. (including My Chiffon Nguyen)
Nature, 2026.

RESEARCH EXPERIENCE

Research Intern (AI Safety)

Jun 2026 – Present

Lida Safety Research Group

Remote

- Contributing to three projects in persona alignment, scalable oversight, and sandbagging
- Mentors: [Linh Le](#) (MILA Postdoc) and [David Williams-King](#) (ERA:AI Research Manager)

Analysis Subteam Lead (Multicultural Reasoning Evaluation)

Apr 2026 – Present

Cohere Labs Open Science Community

Remote

- Co-leading an analysis subteam studying LLM behavior on multicultural riddles, focusing on factual accuracy, hallucination behavior, topic-level patterns, and cross-cultural reasoning failure modes
- Coordinating analysis workflows for a 15-person subteam inside a 100+ researcher open-science initiative, aligning model benchmarking, human evaluation, and synthesis outputs
- Collaborating with Cohere Labs researcher [Julia Kreutzer](#) and cross-functional team leads to produce datasets, posts, and publication-oriented findings

Futurekind Research Fellow (Animals-Aligned Agentic AI Benchmarking)

Apr 2026 – Expected Jul 2026

Electric Sheep, Futurekind Spring Fellowship 2026

Remote

- Evaluating whether implicit values in agentic LLMs are inclusive of animals

[\[Proposal\]](#)

AI Research Mentee (Multilingual Agentic Evaluation)

Feb 2026 – Jun 2026

SEACrowd, SEACrowd 2026 Research Apprenticeship

Remote

- Extended [Tau2-Bench](#), a Tool-Agent-User benchmark for real-world domains, into **SEATauBench** to evaluate low-resource Southeast Asian languages and localized task contexts.
- Designed and ran a translation and localization pipeline across five Southeast Asian languages, preserving executable layers while localizing agent-facing tasks, policies, databases, schemas, and tool surfaces.
- Analyzed multi-turn agent-user-tool trajectories across four experimental settings to diagnose performance degradation and language usage
- Mentors: [Dr. Samuel Cahyawijaya](#) (Cohere Labs) and [Patomporn Payoungkhamdee](#) (VISTEC Thailand)

AI Research Lead Mentee (Scalable Oversight & Chain-of-Thought Monitoring) Oct 2025 – Jun 2026
Algoverse AI Research, AI Research Program Fall 2025 (Mentor: [Yeonwoo Jang \(MATS 8.2\)](#)) *Remote*

- Led a project on weak-to-strong chain-of-thought monitoring for detecting [sandbagging](#)—distinguishing deliberate under-performance from genuine incapability. [\[Code\]](#)
- Designed a reproducible monitoring pipeline with Inspect AI and Inspect Scout; evaluated monitor success across 30+ pairs of reasoning models, including Qwen, DeepSeek, Gemma, and GPT-OSS model families.

Applied Machine Learning Research Intern (3D Object Detection for BMW) Jun 2024 – Aug 2024
Landshut University of Applied Sciences, AI & Mixed Reality Lab *Landshut, Bavaria, Germany*

- Tested PointPillarNet and semantic segmentation to detect and count objects in point cloud data of a parking lot
- Supervisors: [Prof. Eduard Kromer](#) and [Prof. Sandra Eisenreich](#)

TEACHING & MENTORING EXPERIENCE

Curious Cardinals, AI & Data Science, Mentor Nov 2025 – Present

- Mentoring a project about robustness of fact-checking language models under evidence corruption and language shifts

Minerva University, PR51 Programming with Python, Lead Tutor Spring 2025

- Taught 40+ first-year students from 20+ countries in **weekly hands-on programming labs** for 11 weeks, covering Python, OOP, debugging, security, and computing fundamentals
- Analyzed student performance data and tutor surveys across 11 weeks to identify 12 learning bottlenecks, improving engagement metrics by 15% for the next cohort

Minerva University, FA50/FA51 Logic, Probability & Statistics, Lead Teaching Assistant Fall 2022 – Spring 2024

- Supported **150+ students each semester** per semester in formal logic, probability and statistics, algorithmic thinking, and simulation, through weekly office hours
- Provided **formative assessment on 25 semesterly quizzes** for 50 students to correct and shape their learning
- Assisted professors in **grading** three semesterly math and programming assignments

SELECTED PROJECTS

SportConnect: Connect Through Local Sport Events (github.com/mychiffonn/sport-connect) Nov 2025

Scalable full-stack web application connecting users to local recreational sport events, featuring comprehensive event management, event discovery and filtering, and real-time RSVP tracking. Built with TypeScript, React, Express, PostgreSQL, TailwindCSS, DaisyUI, and BetterAuth (for multi-provider OAuth integration).

Mnemonic Generation for Vocabulary Learning (github.com/mychiffonn/mnemonic-gen) Oct 2024 – Mar 2025

- Designed an AI chatbot that generated diverse, memorable, and linguistically-grounded mnemonic devices for learning English and Mandarin Chinese vocabulary
- Utilized chain-of-thought distillation from a teacher model (DeepSeekR1) to instill linguistic chain-of-thought reasoning to a student model (Gemma3-1b) through **supervised fine-tuning (QLoRA)** with unsloth
- Implemented a **Direct Preference Optimization (DPO)** pipeline through trl for preference modeling on Gemma3-1b using 500 human and LLM-annotated preference pairs (on memorability, imageability, and retention rates)

Mini-LLaMA2 PyTorch Implementation (<https://github.com/mychiffonn/cmu-advanced-nlp-minllama>) May 2024

- Implemented a minimalist Llama-2 language model from scratch in PyTorch, including self-attention, Rotary Positional Embeddings (RoPE), RMSNorm, SwiGLU activation functions, and core transformer blocks
- Pretrained an 8-layer, roughly 42M-parameter model on TinyStories and evaluated text completion, zero-shot sentiment, and task-specific finetuning on classification tasks.

Bayesian Modeling of General Practitioner Visits Nov 2023

Modeled general practitioner visit counts across 7 regions and 10 age groups using zero-inflated Poisson and Bayesian hierarchical models in PyMC

OPEN SOURCE CONTRIBUTIONS

[SEACrowd](#) Website Designer & Maintainer Aug 2025 – Present

Web infrastructure and communication for a Southeast Asian AI research community

[Scholar \(Astro Theme for Academics\)](#) Creator & Maintainer Jun 2025 – Present

TECHNICAL SKILLS

- **Programming Languages:** Python, SQL, Bash, TypeScript, JavaScript, R

- **Machine Learning:** PyTorch, unsloth, HuggingFace, scikit-learn, Inspect AI, harbor, LangGraph, langfuse
- **Data and Analytics:** PostgreSQL, ClickHouse, dbt, PySpark, MongoDB, Tableau
- **Web and Product Engineering:** React, FastAPI, Astro, Express.js, Flask, TailwindCSS, shadcn/ui
- **Tools and DevOps:** Git, Docker, Vercel, Railway, LaTeX, Zotero

LANGUAGES

Vietnamese (native), English (fluent/C1), Mandarin Chinese (lower-intermediate/B1/HSK 4)

CERTIFICATES

- **Tiny Aya Expedition**, Cohere Labs (credsverse.com/credentials/0e262642-a662-4d2b-ac4e-8ca35fc2f80c) Mar 2026
- **Advanced Web Development**, CodePath ([drive.google.com/file/d/1n4dHj4TFM8HWIDXMTt9ZGjEXVIpkP-F-/](https://drive.google.com/file/d/1n4dHj4TFM8HWIDXMTt9ZGjEXVIpkP-F-/view))
- **Natural Language Specialization**, deeplearning.ai (coursera.org/verify/specialization/3FJ3W7QJX8GK) Nov 2023
- **Data Science Lab**, World Quant University (credly.com/badges/2e1e6902-aae4-47c4-97e2-0ad9265e5561) Aug 2023
- **Machine Learning Specialization**, deeplearning.ai (coursera.org/verify/specialization/G9898XKB9EAV) Jun 2022

Last Updated: Jun 15, 2026