My (Chiffon) Nguyen

San Francisco, CA, USA | hi@mychiffonn.com | github.com/mychiffonn | mychiffonn.com

Research Interests

Current and future AI that are socially responsible and beneficial for broader world: **AI alignment** (pluralistic alignment, goal misgeneralization) and **AI control** (monitoring, scalable oversight); **multilingual and multicultural AI**; **AI x animals**; **human-AI collaboration** and interaction-centric models.

Education

Minerva University, College of Computational Sciences — Sep 2021 – May 2025
B.Sc in Computational Sciences (Machine Learning and Statistics), GPA: 3.7/4.0 — *San Francisco, CA, USA*

- **Relevant Coursework**: Machine Learning (A), Bayesian Modeling (A), Statistical Modeling and Causal Inference (A), Optimization Methods (A), Probability and Statistics (A-)
- **Global Experience**: Seoul (South Korea), Taipei, Hyderabad (India), Buenos Aires (Argentina), Berlin (Germany)

Selected Publication

P=Preprint, **R**=Under Review, **J**=Journal, **C**=Conference, **W**=Workshop. Full publication list on Google Scholar.

[R1] **SEATauBench: Adapting Tool-Agent-User Evaluation Into Low-Resource Southeast Asian Languages**
My Chiffon Nguyen*, Aulia Adila*, Saksorn Ruangtanusak*, Kittiphat Leesombatwathana*, Vissuta Gunawan Lim*, Patomporn Payoungkhamdee, Samuel Cahyawijaya
Under review at EMNLP, 2026.

[P1] **Anthropogenic Regional Adaptation in Multimodal Vision-Language Model**
Samuel Cahyawijaya, Peerat Limkonchotiwat, Tack Hwa Wong, et al. (including **My Chiffon Nguyen**)
Preprint, 2026.

Research Experience

Reearch Intern (AI Safety) — Jun 2026 – Expected Sep 2026
Lida Safety Research Group — *Remote*

- Contributing to two projects in mechanistic interpretability and sandbagging
- Mentors: Linh Le (MILA Postdoc) and David Williams-King (ERA:AI Research Manager)

Analysis Subteam Lead (Multicultural Reasoning Evaluation) — Apr 2026 – Present
Cohere Labs Open Science Community — *Remote*

- Co-leading an analysis subteam studying LLM behavior on multicultural riddles, focusing on factual accuracy, hallucination behavior, topic-level patterns, and cross-cultural reasoning failure modes
- Coordinating analysis workflows for a 15-person subteam inside a 100+ researcher open-science initiative, aligning model benchmarking, human evaluation, and synthesis outputs

Futurekind Research Fellow (Animals-Aligned Agentic AI Benchmarking) — Apr 2026 – Expected Jul 2026
Electric Sheep, Futurekind Spring Fellowship 2026 — *Remote*

- Evaluating whether implicit values in agentic LLMs are inclusive of animals [Proposal]

AI Research Mentee (Multilingual Agentic Evaluation) — Feb 2026 – Jun 2026
SEACrowd, SEACrowd 2026 Research Apprenticeship — *Remote*

- Extended Tau2-Bench, a Tool-Agent-User benchmark for real-world domains, into **SEATauBench** to evaluate low-resource Southeast Asian languages and localized task contexts.
- Designed and ran a translation and localization pipeline across five Southeast Asian languages, preserving executable layers while localizing agent-facing tasks, policies, databases, schemas, and tool surfaces.
- Analyzed multi-turn agent-user-tool trajectories to diagnose performance degradation and language usage
- Mentors: Dr. Samuel Cahyawijaya (Cohere Labs) and Patomporn Payoungkhamdee (VISTEC Thailand)

AI Research Lead Mentee (Scalable Oversight & Chain-of-Thought Monitoring) — Oct 2025 – Jun 2026
Algoverse AI Research, AI Research Program Fall 2025 (Mentor: Yeonwoo Jang (MATS 8.2)) — *Remote*

- Led a project on weak-to-strong chain-of-thought monitoring for detecting sandbagging–distinguishing deliberate underperformance from genuine incapability. [Code]
- Designed a reproducible monitoring pipeline with Inspect AI and Inspect Scout; evaluated monitor success across 30+ pairs of reasoning models, including Qwen, DeepSeek, Gemma, and GPT-OSS model families.

Applied Machine Learning Research Intern (3D Object Detection for BMW) Jun 2024 – Aug 2024
Landshut University of Applied Sciences, AI & Mixed Reality Lab *Landshut, Bavaria, Germany*

- Tested PointPillarNet and semantic segmentation to detect and count objects in point cloud data of a parking lot
- Supervisors: Prof. Eduard Kromer and Prof. Sandra Eisenreich

Teaching & Mentoring Experience

Curious Cardinals, *AI & Data Science*, Mentor Nov 2025 – Present

- Mentoring a project about robustness of fact-checking language models under evidence corruption and language shifts

Minerva University, *PR51 Programming with Python*, Lead Tutor Spring 2025

- Taught 40+ first-year students from 20+ countries in **weekly hands-on programming labs** for 11 weeks, covering Python, OOP, debugging, security, and computing fundamentals
- Analyzed student performance data and tutor surveys across 11 weeks to identify 12 learning bottlenecks, improving engagement metrics by 15% for the next cohort

Minerva University, *FA50/FA51 Logic, Probability & Statistics*, Lead Teaching Assistant Fall 2022 – Spring 2024

- Supported **150+ students each semester** per semester in formal logic, probability and statistics, algorithmic thinking, and simulation, through weekly office hours
- Provided **formative assessment on 25 semesterly quizzes** for 50 students to correct and shape their learning
- Assisted professors in **grading** three semesterly math and programming assignments

Selected Projects

Mnemonic Generation for Vocabulary Learning (github.com/mychiffonn/mnemonic-gen) Oct 2024 – Mar 2025

- Designed an AI chatbot that generated diverse, memorable, and linguistically-grounded mnemonic devices for learning English and Mandarin Chinese vocabulary
- Utilized chain-of-thought distillation from a teacher model (DeepSeekR1) to instill linguistic chain-of-though reasoning to a student model (Gemma3-1b) through **supervised fine-tuning (QLoRA)** with `unsloth`
- Implemented a **Direct Preference Optimization (DPO)** pipeline through `trl` for preference modeling on Gemma3-1b using 500 human and LLM-annotated preference pairs (on memorability, imageability, and retention rates)

Mini-LLaMA2 PyTorch Implementation (https://github.com/mychiffonn/cmu-advanced-nlp-minllama) May 2024

- Implemented a minimalist Llama-2 language model from scratch in PyTorch, including self-attention, Rotary Positional Embeddings (RoPE), RMSNorm, SwiGLU activation functions, and core transformer blocks
- Pretrained an 8-layer, roughly 42M-parameter model on TinyStories and evaluated text completion, zero-shot sentiment, and task-specific finetuning on classification tasks.

SEACrowd Website & Content (https://seacrowd.org/) Aug 2025 – Present

- Designed website and developed communication for a Southeast Asian AI research community

Technical Skills

- **Programming Languages**: Python, SQL, Bash, TypeScript, JavaScript, R
- **Machine Learning**: PyTorch, unsloth, HuggingFace, scikit-learn, Inspect AI, harbor, LangGraph, langfuse
- **Data and Analytics**: PostgreSQL, ClickHouse, dbt, PySpark, MongoDB, Tableau
- **Web and Product Engineering**: React, FastAPI, Astro, Express.js, Flask, TailwindCSS, shadcn/ui
- **Tools and DevOps**: Git, Docker, Vercel, Railway, LaTeX, Zotero
- **Languages**: Vietnamese (native), English (fluent/C1), Mandarin Chinese (lower-intermediate/B1/HSK 4)

Certificates

- **Tiny Aya Expedition**, Cohere Labs (credsverse.com/credentials/0e262642-a662-4d2b-ac4e-8ca35fc2f80c) Mar 2026
- **Advanced Web Development**, CodePath (drive.google.com/file/d/1n4dHj4TFM8HWlDXMTt9ZGjEXVIpkP-F-)
- **Natural Language Specialization**, deeplearning.ai (coursera.org/verify/specialization/3FJ3W7QJX8GK) Nov 2023
- **Machine Learning Specialization**, deeplearning.ai (coursera.org/verify/specialization/G9898XKB9EAV) Jun 2022