

# My (Chiffon) Nguyen

San Francisco, CA, USA | [hi@mychiffonn.com](mailto:hi@mychiffonn.com) | [github.com/mychiffonn](https://github.com/mychiffonn) | [mychiffonn.com](https://mychiffonn.com)

## RESEARCH INTERESTS

---

Current and future AI that are socially responsible and beneficial for broader world: **AI alignment** (pluralistic alignment, goal misgeneralization) and **AI control** (monitoring, scalable oversight); **multilingual and multicultural AI**; **AI x animals**; **human-AI collaboration** and interaction-centric models.

## EDUCATION

---

Minerva University, College of Computational Sciences

Sep 2021 – May 2025

*B.Sc in Computational Sciences (Machine Learning and Statistics), GPA: 3.7/4.0*

*San Francisco, CA, USA*

- **Relevant Coursework:** Machine Learning (A), Bayesian Modeling (A), Statistical Modeling and Causal Inference (A), Optimization Methods (A), Probability and Statistics (A-)
- **Global Experience:** Seoul (South Korea), Taipei, Hyderabad (India), Buenos Aires (Argentina), Berlin (Germany)

## SELECTED PUBLICATION

---

P=Preprint, R=Under Review, J=Journal, C=Conference, W=Workshop. Full publication list on [Google Scholar](https://scholar.google.com/citations?user=MyChiffon).

[R2] **Position: Synthetic Persona Needs Explicit Grounding and Standardized Reporting**

Yucheng Lu, Xiaoyi Liu, Heming Liu, Hanwen Xing, Shirley Huang, Guanghui Min, Zhengyang Shan, My Chiffon Nguyen, Ishan Gupta, Zonglin Di, Qianfeng Wen, Yifan Simon Liu, Ruoqi Gao, and 31 others

*Under review at COLM Workshop SocialSim, 2026.*

[R1] [SEATauBench: Adapting Tool-Agent-User Evaluation Into Low-Resource Southeast Asian Languages](#)

My Chiffon Nguyen\*, Aulia Adila\*, Saksorn Ruangtanusak\*, Kittiphath Leesombatwathana\*, Vissuta Gunawan Lim\*, Patomporn Payoungkhamdee, Samuel Cahyawijaya

*Under review at EMNLP, 2026.*

[P1] [Anthropogenic Regional Adaptation in Multimodal Vision-Language Model](#)

Samuel Cahyawijaya, Peerat Limkonchotiwat, Tack Hwa Wong, et al. (including My Chiffon Nguyen)

*Preprint, 2026.*

## RESEARCH EXPERIENCE

---

### Research Intern (AI Safety)

Jun 2026 – Expected Sep 2026

Lida Safety Research Group

*Remote*

- Contributing to two projects in mechanistic interpretability and sandbagging
- Mentors: [Linh Le](#) (MILA Postdoc) and [David Williams-King](#) (ERA:AI Research Manager)

### Analysis Subteam Lead (Multicultural Reasoning Evaluation)

Apr 2026 – Present

Cohere Labs Open Science Community

*Remote*

- Co-leading an analysis subteam studying LLM behavior on multicultural riddles, focusing on factual accuracy, hallucination behavior, topic-level patterns, and cross-cultural reasoning failure modes
- Coordinating analysis workflows for a 15-person subteam inside a 100+ researcher open-science initiative, aligning model benchmarking, human evaluation, and synthesis outputs

### AI Research Mentee (Multilingual Agentic Evaluation)

Feb 2026 – Jun 2026

SEACrowd, SEACrowd 2026 Research Apprenticeship

*Remote*

- Extended [Tau2-Bench](#), a Tool-Agent-User benchmark for real-world domains, into **SEATauBench** to evaluate low-resource Southeast Asian languages and localized task contexts.
- Designed and ran a translation and localization pipeline across five Southeast Asian languages, preserving executable layers while localizing agent-facing tasks, policies, databases, schemas, and tool surfaces.
- Analyzed multi-turn agent-user-tool trajectories to diagnose performance degradation and language usage
- Mentors: [Dr. Samuel Cahyawijaya](#) (Cohere Labs) and [Patomporn Payoungkhamdee](#) (VISTEC Thailand)

### AI Research Lead Mentee (Scalable Oversight & Chain-of-Thought Monitoring)

Oct 2025 – Jun 2026

Algoverse AI Research, AI Research Program Fall 2025 (Mentor: [Yeonwoo Jang \(MATS 8.2\)](#))

*Remote*

- Led a project on weak-to-strong chain-of-thought monitoring for detecting [sandbagging](#)—distinguishing deliberate under-performance from genuine incapability. [\[Code\]](#)

- Designed a reproducible monitoring pipeline with Inspect AI and Inspect Scout; evaluated monitor success across 30+ pairs of reasoning models, including Qwen, DeepSeek, Gemma, and GPT-OSS model families.

## Applied Machine Learning Research Intern (3D Object Detection for BMW)

Jun 2024 – Aug 2024

Landshut University of Applied Sciences, AI & Mixed Reality Lab

Landshut, Bavaria, Germany

- Object detection using synthetic point cloud data (BMW), PointPillarNet and semantic segmentation
- Supervisors: [Prof. Eduard Kromer](#) and [Prof. Sandra Eisenreich](#)

## TEACHING & MENTORING EXPERIENCE

---

**Curious Cardinals**, *AI & Data Science*, Mentor

Nov 2025 – Present

- Mentoring a project about robustness of fact-checking language models under evidence corruption and language shifts

**Minerva University**, *PR51 Programming with Python*, Lead Tutor

Spring 2025

- Taught 40+ first-year students from 20+ countries in **weekly hands-on programming labs** for 11 weeks, covering Python, OOP, debugging, security, and computing fundamentals
- Analyzed student performance data and tutor surveys across 11 weeks to identify 12 learning bottlenecks, improving engagement metrics by 15% for the next cohort

**Minerva University**, *FA50/FA51 Logic, Probability & Statistics*, Lead Teaching Assistant

Fall 2022 – Spring 2024

- Supported **150+ students each semester** per semester in formal logic, probability and statistics, algorithmic thinking, and simulation, through weekly office hours
- Provided **formative assessment on 25 semesterly quizzes** for 50 students to correct and shape their learning
- Assisted professors in **grading** three semesterly math and programming assignments

## SELECTED PROJECTS

---

**Mnemonic Generation for Vocabulary Learning** ([github.com/mychiffonn/mnemonic-gen](https://github.com/mychiffonn/mnemonic-gen))

Oct 2024 – Mar 2025

- Designed an AI chatbot that generated diverse, memorable, and linguistically-grounded mnemonic devices for learning English and Mandarin Chinese vocabulary
- Utilized chain-of-thought distillation from a teacher model (DeepSeekR1) to instill linguistic chain-of-thought reasoning to a student model (Gemma3-1b) through **supervised fine-tuning (QLoRA)** with unsloth
- Implemented a **Direct Preference Optimization (DPO)** pipeline through trl for preference modeling on Gemma3-1b using 500 human and LLM-annotated preference pairs (on memorability, imageability, and retention rates)

**Mini-LLaMA2 PyTorch Implementation** (<https://github.com/mychiffonn/cmu-advanced-nlp-minllama>)

May 2024

- Implemented a minimalist Llama-2 language model from scratch in PyTorch, including self-attention, Rotary Positional Embeddings (RoPE), RMSNorm, SwiGLU activation functions, and core transformer blocks
- Pretrained an 8-layer, roughly 42M-parameter model on TinyStories and evaluated text completion, zero-shot sentiment, and task-specific finetuning on classification tasks.

**SEACrowd Website & Content** (<https://seacrowd.org/>)

Aug 2025 – Present

- Designed website and developed communication for a Southeast Asian AI research community

## TECHNICAL SKILLS

---

- **Programming Languages:** Python, SQL, Bash, TypeScript, JavaScript, R
- **Machine Learning:** PyTorch, unsloth, HuggingFace, scikit-learn, Inspect AI, harbor, LangGraph, langfuse
- **Data and Analytics:** PostgreSQL, ClickHouse, dbt, PySpark, MongoDB, Tableau
- **Web and Product Engineering:** React, FastAPI, Astro, Express.js, Flask, TailwindCSS, shadcn/ui
- **Tools and DevOps:** Git, Docker, Vercel, Railway, LaTeX, Zotero
- **Languages:** Vietnamese (native), English (fluent/C1), Mandarin Chinese (lower-intermediate/B1/HSK 4)

## CERTIFICATES

---

- **Tiny Aya Expedition**, Cohere Labs ([credsverse.com/credentials/0e262642-a662-4d2b-ac4e-8ca35fc2f80c](https://credsverse.com/credentials/0e262642-a662-4d2b-ac4e-8ca35fc2f80c)) Mar 2026
- **Advanced Web Development**, CodePath ([drive.google.com/file/d/1n4dHj4TFM8HWIDXMTt9ZGjEXVIpKpF-/view](https://drive.google.com/file/d/1n4dHj4TFM8HWIDXMTt9ZGjEXVIpKpF-/view))
- **Natural Language Specialization**, deeplearning.ai ([coursera.org/verify/specialization/3FJ3W7QJX8GK](https://coursera.org/verify/specialization/3FJ3W7QJX8GK)) Nov 2023
- **Machine Learning Specialization**, deeplearning.ai ([coursera.org/verify/specialization/G9898XKB9EAV](https://coursera.org/verify/specialization/G9898XKB9EAV)) Jun 2022