

# My (Chiffon) Nguyen

Hanoi, Vietnam | [hi@mychiffonn.com](mailto:hi@mychiffonn.com) | [github.com/mychiffonn](https://github.com/mychiffonn) | [mychiffonn.com](https://mychiffonn.com)

## RESEARCH INTERESTS

---

Current and future AI that are socially responsible and beneficial for broader world: **AI alignment** (safety evaluation, pluralistic alignment, goal misgeneralization) and **AI control** (monitoring, scalable oversight); **multilingual and multicultural AI**; cooperative AI for **human-AI collaboration** and **multi-agent systems**.

## SELECTED PUBLICATION & SUBMISSIONS

---

P=Preprint, R=Under Review, J=Journal, C=Conference, W=Workshop. Full publication list on [Google Scholar](https://scholar.google.com/citations?user=MyChiffon).

[W1] **Position: Synthetic Persona Needs Explicit Grounding and Standardized Reporting**

Yucheng Lu, Xiaoyi Liu, Heming Liu, Hanwen Xing, Shirley Huang, Guanghui Min, Zhengyang Shan, **My Chiffon Nguyen**, Ishan Gupta, Zonglin Di, Qianfeng Wen, Yifan Simon Liu, Ruoqi Gao, and 31 others  
*COLM Workshop Social Simulation with LLMs*, 2026.

[R3] **GPS-Bench: A Governance Policy Simulator for Automating Policy Analysis**

Melanie Bui, **My Chiffon Nguyen**, Zachary Schlosser, David Williams-King, Linh Le  
*Under review at AAAI 2027 AI Alignment Track*, 2026.

[R2] **MultiCulturalRiddle: A Multicultural Benchmark of Riddles**

Tianyi Hu\*, Henry Gagnier\*, Vinod Anbalagan\*, **My Chiffon Nguyen\***, Pouya Sadeghi\*, Farah Abdou\*, and 114 others  
*Under review at EMNLP Workshop Multilingual Representation Learning (MRL)*, 2026.

[R1] [SEATauBench: Adapting Tool-Agent-User Evaluation Into Low-Resource Southeast Asian Languages](#)

**My Chiffon Nguyen\***, Aulia Adila\*, Saksorn Ruangtanusak\*, Kittiphat Leesombatwathana\*, Vissuta Gunawan Lim\*, Patomporn Payoungkhamdee, Samuel Cahyawijaya  
*Under review at EMNLP*, 2026.

## RESEARCH EXPERIENCE

---

### Research Intern (AI Governance Simulation)

Jul 2026 – Present

Lida Safety Research Group

*Remote*

- Co-authoring **GPS-Bench**, an evidence-grounded benchmark and simulation framework for evaluating AI-governance policy forecasts against historical records; **under review at AAAI 2027 AI Alignment Track**
- Work on data collection, annotations, and writing
- Mentors: [Dr. Linh Le](#) and [Dr. David Williams-King](#)

### Research Coordinator, Analysis (Multicultural Benchmarking)

Apr 2026 – Present

Cohere Labs Open Science Community

*Remote*

- Co-first authoring **MultiCulturalRiddle**, a culturally-grounded riddle benchmark spanning 61 cultures and 51 languages built through a 100+ researcher participatory community effort; first version **under review at EMNLP 2026 Workshop MRL**
- Collaborating with seven other team leads to build workflows inside the open-science initiative
- Mentors: [Dr. Julia Kreutzer](#) (Cohere Labs), [Károly Boczka](#), and [Prof Akhil Arora](#) (Aarhus University)

### AI Research Mentee (Multilingual Agent Evaluation)

Feb 2026 – Jun 2026

SEACrowd, SEACrowd 2026 Research Apprenticeship

*Remote*

- Co-first authored **SEATauBench**, a benchmark evaluating tool-using conversational agents across four progressive language adaptation scenarios and five low-resource languages, by localizing [Tau2-Bench](#); **under review at EMNLP 2026** [[Code](#)] [[Paper](#)]
- Analyzed multi-turn agent-user-tool trajectories across 63,000+ model outputs to diagnose performance degradation and language usage, surfacing 10 types of tool-calling failures across domains and languages
- Designed and ran a machine translation pipeline that localizes agent-facing contexts while preserving executable and evaluation layers
- Mentors: [Dr. Samuel Cahyawijaya](#) (Cohere Labs) and [Patomporn Payoungkhamdee](#)

### AI Research Lead Mentee (Chain-of-Thought Monitoring)

Oct 2025 – Apr 2026

AlgoVerse AI Research, AI Research Program Fall 2025 (Mentor: [Yeonwoo Jang \(MATS 8.2\)](#))

*San Francisco, CA, US*

- Led a project on weak-to-strong chain-of-thought monitoring for detecting [sandbagging](#)—distinguishing deliberate underperformance from genuine incapability. [[Code](#)] [[Draft](#)]
- Designed a reproducible monitoring pipeline with Inspect AI and Inspect Scout; evaluated monitor success across 30+ pairs of reasoning models on CAIS's WMDP and BigCodeBench

## WORK EXPERIENCE

---

### AI Safety Operations

Jul 2026 – Present

Equistamp

*Remote*

- Contributed to METR's Monitorability project by providing human baselining in transcript safety analysis
- Conducted blue-teaming for Redwood Research's ML Research Sabotage project

## AI Mentor

Curious Cardinals

Dec 2025 – Jul 2026

San Francisco, CA, US

- Mentoring a high schooler on a research project on robustness of fact-checking language models under evidence corruption and language shifts, submitted to \*ACL workshops
- Tutoring selected AI topics, from neural network fundamentals to neolab research trends

## TEACHING & MENTORING EXPERIENCE

---

Minerva University, *PR51 Programming with Python*, Lead Tutor

Jan 2025 – May 2025

- Taught 40+ first-year students from 20+ countries in **weekly hands-on programming labs** for 11 weeks, covering Python, OOP, debugging, security, and computing fundamentals
- Analyzed student performance data and tutor surveys across 11 weeks to identify 12 learning bottlenecks, improving engagement metrics by 15% for the next cohort
- Reviewed and gave feedback on **9 Python programming assignments** for 150+ students

Minerva University, *FA50/FA51 Logic, Probability & Statistics*, Lead Teaching Assistant

Sep 2022 – Apr 2024

- Supported **150+ students** each semester in formal logic, probability and statistics, algorithmic thinking, and simulation, through **70+ weekly office hours**
- Provided **formative assessment on 25 semesterly quizzes** for 50 students to correct and shape their learning
- Assisted professors in **grading** three semesterly math and programming assignments

## SELECTED PROJECTS

---

BenchFlow ([github.com/benchflow-ai/benchflow](https://github.com/benchflow-ai/benchflow))

Jun 2026 – Present

- Contributed **three interactive user simulation** examples to the open-source project on runtime environment for AI agents, modeling multi-turn user-agent interaction in a shared Agent Client Protocol (ACP) sandbox

MyScholar Astro Theme ([github.com/mychiffonn/myscholar](https://github.com/mychiffonn/myscholar))

Nov 2025 – Present

- Designed a theme for academic portfolios and labs in (Astro v7, TypeScript), with publication export, technical writing, and i18n

Multilingual Representations ([labscommunity.cohere.com/blog/2026/inside-tiny-aya-what-an-sae-sees/](https://labscommunity.cohere.com/blog/2026/inside-tiny-aya-what-an-sae-sees/))

Mar 2026

- Investigated whether circuits in multilingual LLM like Tiny Aya are language-specific, script-specific, regional, or universal with Sparse Autoencoders (SAEs) [[HuggingFace](https://huggingface.co)]
- Evaluated generation quality before and after feature steering across 67 languages using LaBSE semantic similarity, script conformity, and LLM-as-a-judge, comparing with Flores-200 and CulturaX datasets

## TECHNICAL SKILLS

---

- **Programming Languages:** Python, SQL, Bash, TypeScript, JavaScript, R
- **Machine Learning:** PyTorch, unsloth, HuggingFace, scikit-learn, Inspect AI, harbor, LangGraph, langfuse
- **Web and Product Engineering:** React, FastAPI, Astro, Next.js, Express.js, Flask, TailwindCSS
- **Tools and DevOps:** Git, Docker, Vercel, Railway, LaTeX, Zotero
- **Languages:** Vietnamese (native), English (fluent/C1), Mandarin Chinese (lower-intermediate/B1/HSK 4)

## EDUCATION

---

Minerva University, College of Computational Sciences

Sep 2021 – May 2025

*B.Sc in Computational Sciences (Machine Learning and Statistics), GPA: 3.7/4.0*

San Francisco, CA, USA

- **Relevant Coursework:** Data Structures & Algorithms, Software Engineering, Machine Learning (A), Bayesian Modeling (A), Statistical Modeling and Causal Inference (A), Optimization Methods (A), Probability and Statistics (A-)
- **Global Experience:** Seoul (South Korea), Taipei, Hyderabad (India), Buenos Aires (Argentina), Berlin (Germany)

## CERTIFICATES

---

• **Technical AI Alignment (ARENA)**, BASE ([arena.baseresearch.org](https://arena.baseresearch.org))

Jul - Sep 2026

• **Natural Language Specialization**, deeplearning.ai ([coursera.org/verify/specialization/3FJ3W7QJX8GK](https://coursera.org/verify/specialization/3FJ3W7QJX8GK))

Nov 2023

• **Machine Learning Specialization**, deeplearning.ai ([coursera.org/verify/specialization/G9898XKB9EAV](https://coursera.org/verify/specialization/G9898XKB9EAV))

Jun 2022

*Last Updated: Aug 20, 2026*