

My (Chiffon) Nguyen

Hanoi, Vietnam | hi@mychiffonn.com | github.com/mychiffonn | mychiffonn.com

RESEARCH INTERESTS

Current and future AI that are socially responsible and beneficial for broader world: **AI alignment** (safety evaluation, pluralistic alignment, goal misgeneralization) and **AI control** (monitoring, scalable oversight); **multilingual and multicultural AI**; cooperative AI for **human-AI collaboration** and **multi-agent systems**.

EDUCATION

Minerva University, College of Computational Sciences

Sep 2021 – May 2025

B.Sc in Computational Sciences (Machine Learning and Statistics), GPA: 3.7/4.0

San Francisco, CA, USA

- **Relevant Coursework:** Data Structures & Algorithms, Software Engineering, Machine Learning (A), Bayesian Modeling (A), Statistical Modeling and Causal Inference (A), Optimization Methods (A), Probability and Statistics (A-)
- **Global Experience:** Seoul (South Korea), Taipei, Hyderabad (India), Buenos Aires (Argentina), Berlin (Germany)

SELECTED PUBLICATION & SUBMISSIONS

P=Preprint, R=Under Review, J=Journal, C=Conference, W=Workshop. Full publication list on [Google Scholar](https://scholar.google.com/citations?user=MyChiffon).

[R2] CulturalRiddles: A Multicultural Benchmark of Riddles

My Chiffon Nguyen*, Tianyi Hu*, Henry Gagnier*, Vinod Anbalagan*, Pouya Sadeghi*, and 149 others
EMNLP Workshop Multilingual Representation Learning (MRL). Full paper under review, 2026.

[R1] GPS-Bench: A Governance Policy Simulator for Automating Policy Analysis

Linh Le, Melanie Bui, My Chiffon Nguyen, Zachary Schlosser, David Williams-King
Under review, 2026.

[C1] SEATauBench: Adapting Tool-Agent-User Evaluation Into Low-Resource Southeast Asian Languages

My Chiffon Nguyen*, Aulia Adila*, Saksorn Ruangtanusak*, Kittiphat Leesombatwathana*, Vissuta Gunawan Lim*, Patomporn Payoungkhamdee, Samuel Cahyawijaya
EMNLP Findings, 2026.

[W1] Position: Synthetic Persona Needs Explicit Grounding and Standardized Reporting

Yucheng Lu, Xiaoyi Liu, Heming Liu, Hanwen Xing, Shirley Huang, Guanghui Min, Zhengyang Shan, My Chiffon Nguyen, Ishan Gupta, Zonglin Di, Qianfeng Wen, Yifan Simon Liu, Ruoqi Gao, and 31 others
COLM Workshop Social Simulation with LLMs, 2026.

RESEARCH EXPERIENCE

SPAR Research Mentee (AI Policy Simulation)

Jul 2026 – Present

Supervised Program for Alignment Research (SPAR) Fall 2026

Remote

- Co-authoring **GPS-Bench**, an evidence-grounded benchmark and simulation framework for evaluating AI-governance policy forecasts against historical records; driving contributions on large-scale data collection, graph building, and writing
- Mentors: [Dr. Linh Le](#) (Lida Safety) and [Dr. David Williams-King](#) (Lida Safety, ERA:AI)

Project Co-Lead (Multicultural Reasoning Evaluation)

Apr 2026 – Oct 2026

Cohere Labs Community

Remote

- Co-first authoring and co-leading **MultiCulturalRiddle**, a culturally-grounded riddle benchmark spanning 61 cultures and 51 languages built through a 100+ researcher participatory community effort
- Leading analyses of 146,400 model responses with human annotations and LLM judges; writing manuscript
- Advisors: [Dr. Julia Kreutzer](#) (Cohere Labs), [Károly Boczka](#), and [Prof Akhil Arora](#) (Aarhus University)

SEACrowd Research Mentee (Multilingual Agent Evaluation)

Feb 2026 – Jun 2026

SEACrowd, SEACrowd 2026 Research Apprenticeship

Remote

- Co-first authored **SEATauBench**, a benchmark evaluating tool-using conversational agents across four progressive language adaptation scenarios and five low-resource languages, by localizing [Tau2-Bench](#); **EMNLP Findings 2026** [\[Code\]](#) [\[Paper\]](#)
- Designed machine translation pipeline for agent-facing contexts while preserving executable and scoring layers, and ran an experiment (L2 Domain); conducted four-step review of translated trajectories, including managing human study for five language quality audits
- Analyzed 63000+ multi-turn agent-user-tool trajectories to diagnose performance degradation and language usage
- Mentors: [Dr. Samuel Cahyawijaya](#) (Cohere Labs) and [Dr. Patomporn Payoungkhamdee](#)

AI Research Lead Mentee (Chain-of-Thought Monitoring)

Oct 2025 – Apr 2026

AlgoVerse AI Research, AI Research Program Fall 2025 (Mentor: [Yeonwoo Jang](#))

San Francisco, CA, US

- Led a project on weak-to-strong chain-of-thought monitoring for detecting [sandbagging](#)—distinguishing deliberate underperformance from genuine incapability. [\[Code\]](#) [\[Draft\]](#)
- Designed a reproducible monitoring pipeline with Inspect AI and Inspect Scout; evaluated monitor success across 30+ pairs of reasoning models on CAIS's WMDP and BigCodeBench

WORK EXPERIENCE

Research Engineer (AI Safety, Contract)

Jul 2026 – Present

Equistamp

Remote

- Contributed to METR's Monitorability project by providing human baselining in transcript safety analysis
- Conducted blue-teaming for Redwood Research's ML Research Sabotage project

AI Mentor (Contract)

Dec 2025 – Jul 2026

Curious Cardinals

San Francisco, CA, US

- Mentoring a high schooler on a research project on robustness of fact-checking language models under evidence corruption and language shifts, submitted to *ACL workshops
- Tutoring selected AI topics, from neural network fundamentals to neolab research trends

TEACHING & MENTORING EXPERIENCE

Minerva University, PR51 Programming with Python, Lead Tutor

Jan 2025 – May 2025

- Taught 40+ first-year students from 20+ countries in **weekly hands-on programming labs** for 11 weeks, covering Python, OOP, debugging, security, and computing fundamentals
- Analyzed student performance data and tutor surveys across 11 weeks to identify 12 learning bottlenecks, improving engagement metrics by 15% for the next cohort
- Reviewed and gave feedback on **9 Python programming assignments** for 150+ students

Minerva University, FA50/FA51 Logic, Probability & Statistics, Lead Teaching Assistant

Sep 2022 – Apr 2024

- Supported **150+ students** each semester in formal logic, probability and statistics, algorithmic thinking, and simulation, through **70+ weekly office hours**
- Provided **formative assessment on 25 semesterly quizzes** for 50 students to correct and shape their learning
- Assisted professors in **grading** three semesterly math and programming assignments

SELECTED PROJECTS

MyScholar Astro Theme (github.com/mychiffonn/myscholar)

Nov 2025 – Present

- Designed a theme for academic portfolios and labs in (Astro v7, TypeScript), with publication export, technical writing, and i18n

Multilingual Representations (labscommunity.cohere.com/blog/2026/inside-tiny-aya-what-an-sae-sees/)

Mar 2026

- Investigated whether circuits in multilingual LLM like Tiny Aya are language-specific, script-specific, regional, or universal with Sparse Autoencoders (SAEs) [[HuggingFace](#)]
- Evaluated generation quality before and after feature steering across 67 languages using LaBSE semantic similarity, script conformity, and LLM-as-a-judge, comparing with Flores-200 and CulturaX datasets

TECHNICAL SKILLS

- **Programming Languages:** Python, SQL, Bash, TypeScript, JavaScript, R
- **Machine Learning:** PyTorch, unsloth, HuggingFace, scikit-learn, Inspect AI, harbor, LangGraph, langfuse
- **Web and Product Engineering:** React, FastAPI, Astro, Next.js, Express.js, Flask, TailwindCSS
- **Tools and DevOps:** Git, Docker, Vercel, Railway, LaTeX, Zotero
- **Languages:** Vietnamese (native), English (fluent/C1), Mandarin Chinese (lower-intermediate/B1/HSK 4)

CERTIFICATES

- **Self-study:** [Technical AI Alignment](#) (ARENA, 2025)
- **Natural Language Specialization**, deeplearning.ai (coursera.org/verify/specialization/3FJ3W7QJX8GK)
- **Machine Learning Specialization**, deeplearning.ai (coursera.org/verify/specialization/G9898XKB9EAV)

Nov 2023

Jun 2022

Last Updated: Oct 02, 2026