

My (Chiffon) Nguyen

San Francisco, CA, USA | hi@mychiffonn.com | github.com/mychiffonn | linkedin.com/in/mychiffonn | mychiffonn.com

EDUCATION

Minerva University, College of Computational Sciences

Sep 2021 – May 2025

B.Sc. in Computational Sciences (Machine Learning & Statistics), GPA: 3.7/4.0

San Francisco, CA, USA

- **Coursework:** Data Structures & Algorithms, Software Engineering, Machine Learning, Probability & Statistics, Statistical Modeling & Causal Inference, Optimization Methods
- **Teaching:** Lead Tutor (Python) and Lead TA (Logic, Probability & Statistics) for five semesters, taught and mentored **300+** students in programming, debugging, and algorithmic thinking

TECHNICAL SKILLS

- **Languages:** Python, TypeScript, JavaScript, SQL, Bash, R
- **AI / ML:** AI evaluation, prompt engineering, agents & tool-calling, RAG (LangGraph, LlamaIndex), fine-tuning (unsloth), PyTorch, Hugging Face, trl, Inspect AI
- **Backend & Data:** FastAPI, Flask, Express.js, REST APIs, PostgreSQL, SQLite, ClickHouse, Prisma
- **Frontend & Product:** React, Astro, Next.js, TailwindCSS, shadcn/ui
- **DevOps:** Git, Docker, CI/CD, Vercel, Railway

EXPERIENCE

Research Engineer | LLM Evaluation & Monitoring

Oct 2025 – Jun 2026

Algoverse AI Research Program

Remote

- Built reproducible **LLM evaluation pipelines** with Inspect AI to run weak-to-strong chain-of-thought monitoring across **30+ reasoning-model pairs** on **WMDP** and **SWE-Bench** benchmarks [\[Code\]](#)
- Designed **LLM-as-a-judge** classifiers that extract and score reasoning traces to detect **sandbagging** (deliberate model underperformance), improving identification of deceptive-reasoning patterns

AI Research Fellow | Agents, Multilingual Evaluation, Machine Translation

Feb 2026 – Jun 2026

SEACrowd 2026 Research Apprenticeship

Remote

- Extended [\$\tau^2\$ -Bench](#) into **SEATauBench**, a benchmark evaluating tool-using **conversational agents** across five low-resource Southeast Asian languages and **278** localized tasks; under review at **EMNLP 2026**
- Engineered a localization pipeline that preserved executable layers while translating tasks, policies, and schemas, enabling reproducible multilingual agent evaluations
- Analyzed **63,000+** multi-turn agent–user–tool trajectories to diagnose tool-calling failure modes (using LLM judge) and language-use degradation across domains and turns

OPEN SOURCE CONTRIBUTIONS

- **Astro Scholar** — creating a static theme for academic websites (Astro, TypeScript, MDX) [\[GitHub\]](#)
- **BenchFlow** — maintaining a Python library for managing runtime environment for AI agents (280+ stars) [\[GitHub\]](#)

SELECTED PROJECTS

SignFast github.com/mychiffonn/signfast — Signer-first petition platform with **personalized ranking**, LLM-generated summaries, and one-click bulk signing; built with Next.js, TypeScript, Prisma, SQLite, TailwindCSS, and Vitest

Ask My Second Brain github.com/mychiffonn/knowledge-rag — A **retrieval-augmented generation (RAG)** system that answers natural-language questions over personal Markdown/Obsidian notes and PDFs; chunked and embedded documents, indexed with **LlamaIndex**, and responses with **OpenAI API**

Mnemonics for Vocabulary Learning github.com/mychiffonn/mnemonic-gen — LLM assistant that generates linguistically grounded mnemonics for English and Mandarin vocabulary; **fine-tuned** Gemma-3-1B (unsloth) on DeepSeek-R1–**distilled** reasoning traces and a **DPO** preference-optimization pipeline (trl) over 500 annotated preference pairs

PUBLICATION PREPRINTS

[P2] [SEATauBench: Adapting Tool-Agent-User Evaluation Into Low-Resource Southeast Asian Languages](#)

My Chiffon Nguyen*, Aulia Adila*, Saksorn Ruangtanusak*, Kittiphath Leesombatwathana*, Vissuta Gunawan Lim*, Patomporn Payoungkhamdee, Samuel Cahyawijaya

Under review at EMNLP, 2026.

[P1] [Anthropogenic Regional Adaptation in Multimodal Vision-Language Model](#)

Samuel Cahyawijaya, Peerat Limkonchotiwat, Tack Hwa Wong, et al. (including My Chiffon Nguyen)

Preprint, 2026.