Chiffon (My) Nguyen

San Francisco, CA, USA | hi@mychiffonn.com | github.com/mychiffonn | linkedin.com/in/mychiffonn | mychiffonn.com

Technical Skills

- **Languages**: Python, TypeScript, JavaScript, SQL, Bash, R
- **AI / ML**: AI evaluation, prompt engineering, agents & tool-calling, RAG (LangGraph, LlamaIndex), fine-tuning, unsloth, PyTorch, Hugging Face, trl, Inspect AI
- **Backend & Data**: FastAPI, Flask, Express.js, REST APIs, PostgreSQL, SQLite, ClickHouse, Prisma
- **Frontend & Product**: React, Astro, Next.js, TailwindCSS, shadcn/ui
- **DevOps**: Git, Docker, CI/CD, Vercel, Railway

Experience

AI Research Engineer | Agents, Multilingual Evaluation, Machine Translation — Feb 2026 – Jun 2026
SEACrowd, SEACrowd 2026 Research Apprenticeship — *Remote*

- Co-authored **SEATauBench**, a benchmark evaluating tool-using **conversational agents** across five low-resource Southeast Asian languages and **214 agent tasks**, by extending Tau2-Bench; under review at **EMNLP 2026** [Paper]
- Engineered a localization pipeline that preserved executable layers while translating tasks, policies, and schemas, enabling reproducible multilingual agent evaluations, reducing critical runtime errors by **80%**
- Built an LLM-as-a-judge evaluation pipeline to analyze **63,000+** multi-turn trajectories at scale, surfacing 10 types of tool-calling failures across domains and languages

Research Engineer | LLM Evaluation & Monitoring — Oct 2025 – Jun 2026
Algoverse AI Research — *Remote*

- Built reproducible LLM evaluation pipelines with Inspect AI to run weak-to-strong chain-of-thought monitoring across **30+ reasoning-model pairs** on CAIS's **WMDP** and OpenAI's **SWE-Bench Verified** benchmarks [Code]
- Designed LLM-as-a-judge classifiers that extract and score **34000+ reasoning and model outputs** to detect **sandbagging** (deliberate model underperformance), improving identification of deceptive reasoning patterns

Lead Teaching Assistant | Logic, Probability & Statistics; Python Programming — Sep 2022 – May 2025
Minerva University, Academics Team — *San Francisco, CA*

- Selected as **Lead TA for 5 semesters**, running **70+ office hours** to debug Python code and teach algorithmic thinking, probability, statistics, and simulation to 150+ students/semester
- Graded **25+ quizzes per semester** for 50 students, identifying recurring conceptual and code-level gaps
- Reviewed **3 Python programming assignments per semester** for 150+ students, giving feedback on correctness, efficiency, and code style

Selected Projects

Benchflow: Python runtime environment for AI agents [285 stars] — [github.com/benchflow-ai/benchflow]

- Contributed **three interactive user simulation** examples, modeling multi-turn user–agent interaction in a shared Agent Client Protocol (ACP) sandbox

SignFast: Signer-first petition platform, **Next.js, Prisma, SQLite, TailwindCSS** — [github.com/mychiffonn/signfast]

- Implemented personalized ranking to surface relevant petitions and LLM-generated summaries for quick comprehension
- Built one-click bulk signing flow to reduce signer friction across multiple petitions

Mnemonics for Vocabulary Learning: unsloth, trl, HuggingFace, PyTorch — (github.com/mychiffonn/mnemonic-gen)

- Built an LLM assistant that generates memory devices for English and Mandarin vocabulary; **fine-tuned** Gemma-3-1B on DeepSeek-R1–**distilled** reasoning traces and ran **DPO** preference optimization on 500 preference pairs

Publications

1. **SEATauBench**: Adapting Tool-Agent-User Evaluation into Low-Resource SEA Languages. **My Chiffon Nguyen** et al., Under review at EMNLP 2026 — arxiv.org/abs/2606.28715
2. Anthropogenic Regional Adaptation in Multimodal Vision-Language Models. Cahyawijaya, Limkonchotiwat, ..., **My Chiffon Nguyen**, et al., Preprint 2026 — arxiv.org/abs/2604.11490

Education

Minerva University, College of Computational Sciences — Sep 2021 – May 2025
B.Sc. in Computational Sciences (Machine Learning & Statistics), GPA: 3.7/4.0 — *San Francisco, CA, USA*

- **Coursework**: Data Structures & Algorithms, Software Engineering, Machine Learning, Probability & Statistics, Statistical Modeling & Causal Inference, Optimization Methods